

tm.plugin.sentiment

Online Sentiment Analysis using R

Mario Annau

Statistical Natural Language Processing
10.12.2010

Table of Contents

- 1 Motivation
- 2 Related work
- 3 The tm library
- 4 tm.plugin.sentiment
- 5 Examples


George just bought a share

...and wants to track its news coverage on the internet



Google News

Search for Microsoft



Advanced news search

News


Top Stories
More sections ▾

All news
Images
Blogs

All recent news
Past hour
Past day
Past week
Past month
2010
2009
2007
2006
1999
Archives

Sorted by relevance
Sorted by date


Results 1 – 10 of about 1,256 for **microsoft msft**. (0.14 seconds)




MICROSOFT TO PLEASE GAMERS ON BLACK FRIDAY WITH LEADING TITLE DISCOUNTS (MSFT) ☆

Zacks.com - 1 hour ago

Nov 23, 2010 (SmartTrend News Watch via COMTEX) -- Officials from **Microsoft Corp.** (NASDAQ: **MSFT**) announced earlier this week that it will give customers Xbox ...

 **MSFT** · **PINK:CMTX**



Today's Novell Deal Helps Microsoft Continue Linux Fight (NOVL, MSFT) ☆

San Francisco Chronicle · [Matt Rosoff](#) · Nov 22, 2010


But more interestingly, a holding company led by **Microsoft** is paying Novell a separate \$450 million cash payment for the rights to more than 800 patents. ...

[Private equity firms to buy Novell](#) Investors' Business Daily

[Attachmate, Diamondback, Hot Trends](#) TheStreet.com

[Novell Agrees To Private-Equity Takeover For \\$2.2B](#) Wall Street Journal


[Xconomy](#)


[all 685 news articles »](#)  **NOVL** · **MSFT**

Market Updates: Hewlett-Packard (NYSE:HPQ), Medtronic (NYSE:MDT), Microsoft ... ☆

Julia Group - 4 hours ago

—MarketWatch Investors that helped companies from **Microsoft** (NASDAQ: **MSFT**) to Wal-Mart Stores (NYSE:WMT) sell bonds at record-low borrowing costs are being ...

 **HPQ** · **MDT** · **WMT**



Google (NASDAQ: GOOG) Faces Challenge from Microsoft (NASDAQ: MSFT) and Facebook ☆


American Consumer News - 7 hours ago

Microsoft (NASDAQ: **MSFT**) and Facebook are allies. They are joining as a competitor against the supremacy of Google (NASDAQ: GOOG). ...

[Google brings MS Office Docs to Google Cloud Connect](#) FierceContentManagement

[Google Hangs a Long Tail on Apps](#) The Money Times

ZDNet UK



4 / 39

Google News Facts

- Biggest News Aggregator around with about **1 billion clicks** per month
- Over **25000** registered **publishers** worldwide
- **40** different regional **editions**
- About **4500** publishers providing **english content** alone
- Google Finance News covers an estimated number of **7000** (mostly unique) Microsoft **articles per year**, or about **20 articles per day**

Who has time to read that?



Motivation for **tm.plugin.sentiment**

- Retrieve content from news sources which are preferably free-of-charge
- Extract the main content from news pages if necessary
- Build up corpus from content which also includes time tags
- Create time series representing the sentiment of news flow over time

We want to extract the news from various sources. . .



Microsoft Tag Cloud ¹



... and create some decent sentiment time series




Table of Contents

1 Motivation

2 Related work

3 The tm library

4 tm.plugin.sentiment

5 Examples

Bo Pang and Lillian Lee

- Provide probably best research overview resource for sentiment analysis with their paper

[Pang Lee, 2008]

Opinion Mining and Sentiment Analysis

*Foundations and Trends in Information Retrieval 2(1-2) ,pp.
1-135 ,2008*

- Their own research focuses on sentiment analysis of online reviews
- Analyzed movie and online product reviews

Paul Tetlock

- Sentiment analysis of popular Wall Street journal column “Abreast of the Market”

[Tetlock, 2007]

Giving Content to Investor Sentiment: The Role of Media in the Stock Market

Journal of Finance 62, 1139-1168, 2007

- Use bag-of-words model and dictionary from the General Inquirer
- Negative news sentiment category (total: 77) has most predictive power for market prices like the Dow Jones Industrial Avg. (DOW 30)

Lydia/Textmap

- Large Scale Sentiment Analysis Project at Stony Brook University, under supervision of Steven Skiena
- Daily analysis of over 1000 english and foreign language online news sources, blogs, rss feeds, historical archives, and other sources
- Online Analysis possible at <http://www.textmap.com/>
- Some technical details:
 - Use (adapted) wget for web spidering
 - Currently stores over one terrabyte of text
 - Use Hadoop Map/Reduce for already tagged text using Amazon's cloud service

Commercial Products

- Thomson One
- Dow Jones News Analytics
- Raven Pack [ravenpack.com] *Bayes training, vector classification, word/phrase lists, pattern detection and market response-based analysis are just a few techniques RavenPack deploys in conducting news sentiment analysis.*

Related Open Source Products/Projects

- Rapidminer with Web Extension
- Python with NLTK

Rapidminer with Web Extension

- Rapidminer based on academic project YALE (TU Dortmund)
- Semi-commercial product
- Workflow oriented
- Implemented in JAVA
- Extensions for Textprocessing and Webmining, even R
- Features include web retrieval, content extraction and bag-of-words analysis

Rapidminer with Web Extension (2)

- Pros:
 - Large set of operators (esp. through WEKA)
 - Workflow orientation for transparency
- Cons:
 - Customization, changes often have to take place in JAVA code

Python with NLTK (1)

- Python is a quite popular scripting language
- Supported by a vast amount of libraries, e.g. the numpy, scipy matplotlib combination for fast numeric computations
- NLTK is a very popular library for text mining researchers
- User can plug together his own sentiment analysis library using packages like NLTK, feedparser, simplejson, etc.

Python with NLTK (2)

- Pros:
 - Excellent scripting language, also interactive
 - Quality of libraries (especially for our purposes: NLTK)
 - Can easily integrated with GUI's (e.g. PyQt)
- Cons:
 - Compared to R: Worse support for stat. functions

Table of Contents

- 1 Motivation
- 2 Related work
- 3 The tm library**
- 4 tm.plugin.sentiment
- 5 Examples

tm – Short Intro

- Developed by Feinerer as part of his dissertation in 2008
- Provides basic data structures for storing large text corpora
- Abstracts data sources and readers for input
- Includes functions for preprocessing, annotation
- Connectors to various open source libraries like **openNLP**, **KEA**

Basic Data Structures

- Text documents, storing texts and individual text meta information
- Corpora, storing collections of Text Documents and meta data
 - DMeta(): Store Classification results for each text document
 - CMeta(): General Corpus Information like creation date
- Possibility for database storage if corpus does not fit into memory ('PCorpus' vs. 'VCorpus').

Source-Reader Concept

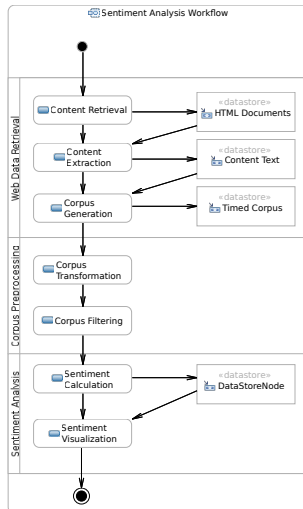
- Abstract data source and reader functionality
- **Source:** Specifies how to access elements and move forward
- **Reader:** Specifies, how content can be extracted and put in a Text Document data structure.
- Call of function `Corpus()` lets you freely specify Source and Reader: `r <- Corpus(DirSource(reut21578), readerControl = list(reader = readReut21578XMLasPlain))`

Table of Contents

- 1 Motivation
- 2 Related work
- 3 The tm library
- 4 tm.plugin.sentiment**
- 5 Examples

Overview

- Web Data Retrieval
- Preprocessing
- Sentiment Analysis



Data retrieval

- Feeds provide meta data of news content about any topic
- Actual content usually resides on external web pages
- Feeds contain urls to content pages

Data retrieval (2)

Therefore a 2-step procedure is necessary:

- 1 Download meta data feeds: `getFeed()`
- 2 Download content sites (optional): `getURLPart()`



Sources Available

For ease of use many different news sources have already been implemented in **tm.plugin.sentiment**:

- News Sources from RSS Feeds like Google News, Yahoo News
- APIs (free registrations required) like Reuters Spotlight, NY Times, Yahoo BOSS, Bing



Content Extraction


- News Sites often contain side bars, headers and ads
- We are only interested in the actual news story
- Heuristics are needed in order to get rid of the 'fluff'
- Best source for a complete overview of HTML extraction techniques:

[Gotttron, 2008]

Content Extraction: Identifying the Main Content in HTML Documents

Johannes Gutenberg-University, Mainz, 2008

Content Extraction (2)



Source: AI Depot

Content Extraction (3)

Implemented in function `extractContentDOM()`:

- 1 Examine each HTML subnode from top to bottom
- 2 If $\text{textlength}/\text{totallength} < \text{threshold}$ then drill down
- 3 Select text from subnode with longest textlength

Idea stems AI Depot and Jinliang Song's ExtMainText Python code

Preprocessing

After Content retrieval/extraction some standard tm preprocessing steps may be required depending on source/data quality.

- Use of tm transformation functions like `tolower()`, `removePunctuation()`, etc.
- Extract “interesting” part of text content using e.g. `getRelevant()`. Useful for sites like [this](#).



Sentiment Calculation

As a first attempt `tm.plugin.sentiment` uses bags-of-words model for sentiment calculation. Therefore the following ingredients are needed:

- Document Term Matrix
- Dictionary of Sentiment-laden words like *good*, *happy*, *loose* or *bankrupt*. Available from General Inquirer, NTU Sentiment Dictionary, OpinionFinder's Subjectivity Lexicon or SentiWordnet

In order to build time series from sentiment scores each document needs to be timestamped.

Sentiment Indicators²

$$polarity = \frac{p - n}{p + n} \quad (1)$$

$$subjectivity = \frac{n + p}{N} \quad (2)$$

$$pos_refs_per_ref = \frac{p}{N} \quad (3)$$

$$neg_refs_per_ref = \frac{n}{N} \quad (4)$$

$$senti_diffs_per_ref = \frac{p - n}{N} \quad (5)$$

²taken from the Lydia/Textmap project



Sentiment Visualization




Table of Contents

- 1 Motivation
- 2 Related work
- 3 The tm library
- 4 tm.plugin.sentiment
- 5 Examples**

Examples

- Google Finance News — Microsoft (MSFT)
- Yahoo BOSS — Alcoa (AA)
- Reuters Spotlight — Gold Market (GLD)